

CHAPTER 3: PROCESSES AND PROCESS SCHEDULING

Dr\ Eman Monir

CHAPTER 3: PROCESSES

- PROCESS CONCEPT
- PROCESS SCHEDULING
- OPERATIONS ON PROCESSES
- INTERPROCESS COMMUNICATION
- EXAMPLES OF IPC SYSTEMS
- COMMUNICATION IN CLIENT-SERVER SYSTEMS

OBJECTIVES

- TO INTRODUCE THE NOTION OF A PROCESS -- A PROGRAM IN EXECUTION, WHICH FORMS THE BASIS OF ALL COMPUTATION
- TO DESCRIBE THE VARIOUS FEATURES OF PROCESSES, INCLUDING SCHEDULING, CREATION AND TERMINATION, AND COMMUNICATION
- TO EXPLORE INTERPROCESS COMMUNICATION USING SHARED MEMORY AND MESSAGE PASSING
- TO DESCRIBE COMMUNICATION IN CLIENT-SERVER SYSTEMS


PROCESS CONCEPT

- AN OPERATING SYSTEM EXECUTES A VARIETY OF PROGRAMS:
 - BATCH SYSTEM – **JOBS**
 - TIME-SHARED SYSTEMS – **USER PROGRAMS** OR **TASKS**
- TEXTBOOK USES THE TERMS *JOB* AND *PROCESS* ALMOST INTERCHANGEABLY
- **PROCESS** – A PROGRAM IN EXECUTION; PROCESS EXECUTION MUST PROGRESS IN SEQUENTIAL FASHION
- MULTIPLE PARTS
 - THE PROGRAM CODE, ALSO CALLED **TEXT SECTION**
 - CURRENT ACTIVITY INCLUDING **PROGRAM COUNTER**, PROCESSOR REGISTERS
 - **STACK** CONTAINING TEMPORARY DATA
 - FUNCTION PARAMETERS, RETURN ADDRESSES, LOCAL VARIABLES
 - **DATA SECTION** CONTAINING GLOBAL VARIABLES
 - **HEAP** CONTAINING MEMORY DYNAMICALLY ALLOCATED DURING RUN TIME

PROCESS CONCEPT (CONT.)

- PROGRAM IS *PASSIVE* ENTITY STORED ON DISK (**EXECUTABLE FILE**), PROCESS IS *ACTIVE*
 - PROGRAM BECOMES PROCESS WHEN EXECUTABLE FILE LOADED INTO MEMORY
- EXECUTION OF PROGRAM STARTED VIA GUI MOUSE CLICKS, COMMAND LINE ENTRY OF ITS NAME, ETC
- ONE PROGRAM CAN BE SEVERAL PROCESSES
 - CONSIDER MULTIPLE USERS EXECUTING THE SAME PROGRAM

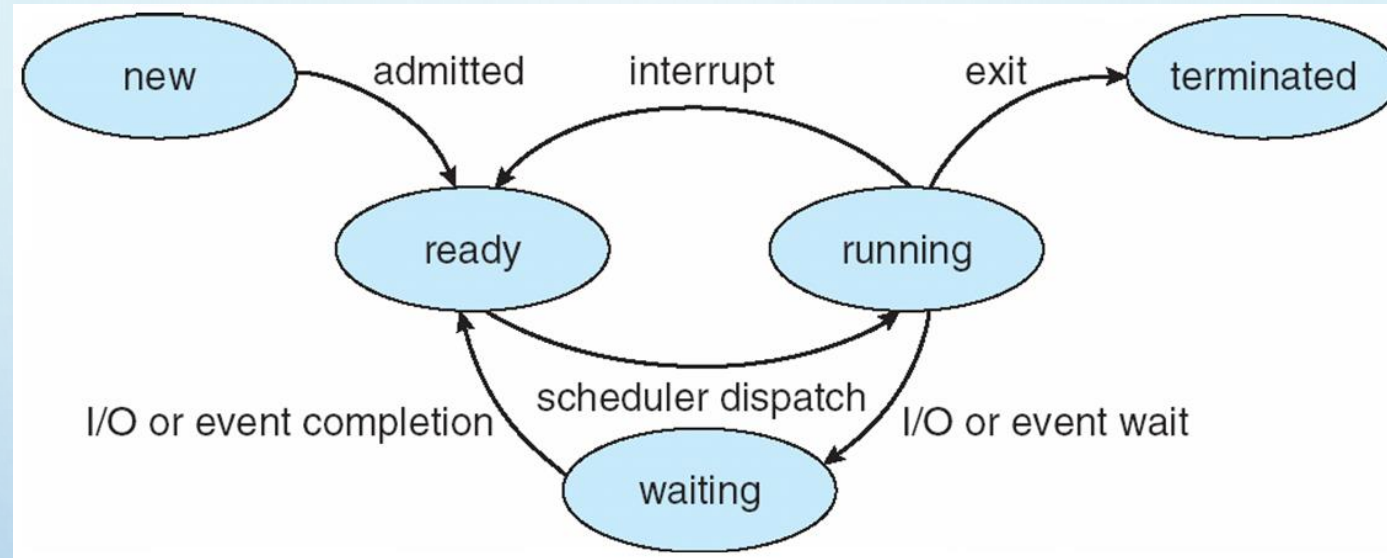
PROCESS IN MEMORY



PROCESS STATE

- AS A PROCESS EXECUTES, IT CHANGES **STATE**
 - **NEW:** THE PROCESS IS BEING CREATED
 - **RUNNING:** INSTRUCTIONS ARE BEING EXECUTED
 - **WAITING:** THE PROCESS IS WAITING FOR SOME EVENT TO OCCUR
 - **READY:** THE PROCESS IS WAITING TO BE ASSIGNED TO A PROCESSOR
 - **TERMINATED:** THE PROCESS HAS FINISHED EXECUTION

DIAGRAM OF PROCESS STATE

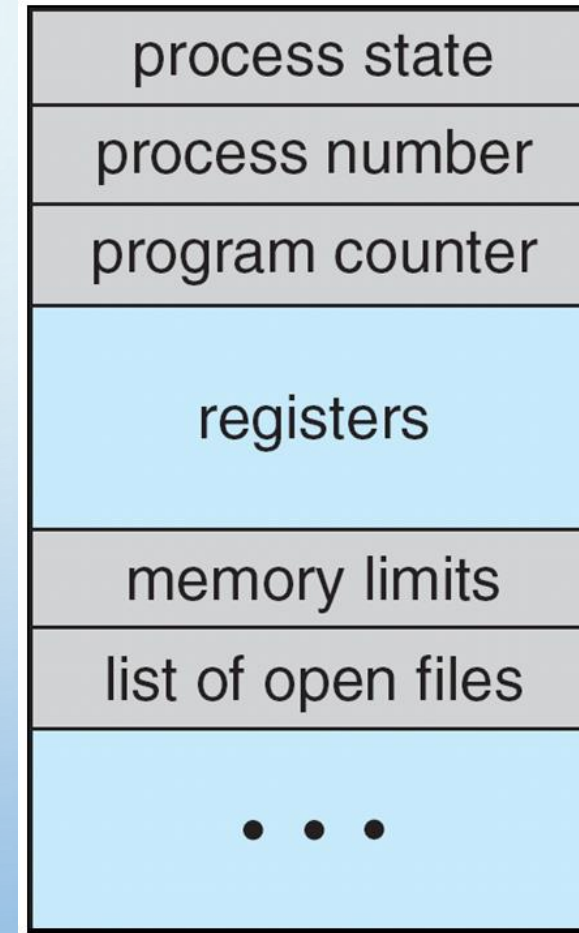


PROCESS CONTROL BLOCK (PCB)

INFORMATION ASSOCIATED WITH EACH PROCESS

(ALSO CALLED **TASK CONTROL BLOCK**)

- PROCESS STATE – RUNNING, WAITING, ETC
- PROGRAM COUNTER – LOCATION OF INSTRUCTION TO NEXT EXECUTE
- CPU REGISTERS – CONTENTS OF ALL PROCESS-CENTRIC REGISTERS
- CPU SCHEDULING INFORMATION– PRIORITIES, SCHEDULING QUEUE POINTERS
- MEMORY-MANAGEMENT INFORMATION – MEMORY ALLOCATED TO THE PROCESS
- ACCOUNTING INFORMATION – CPU USED, CLOCK TIME ELAPSED SINCE START, TIME LIMITS
- I/O STATUS INFORMATION – I/O DEVICES ALLOCATED TO PROCESS, LIST OF OPEN FILES



THREADS

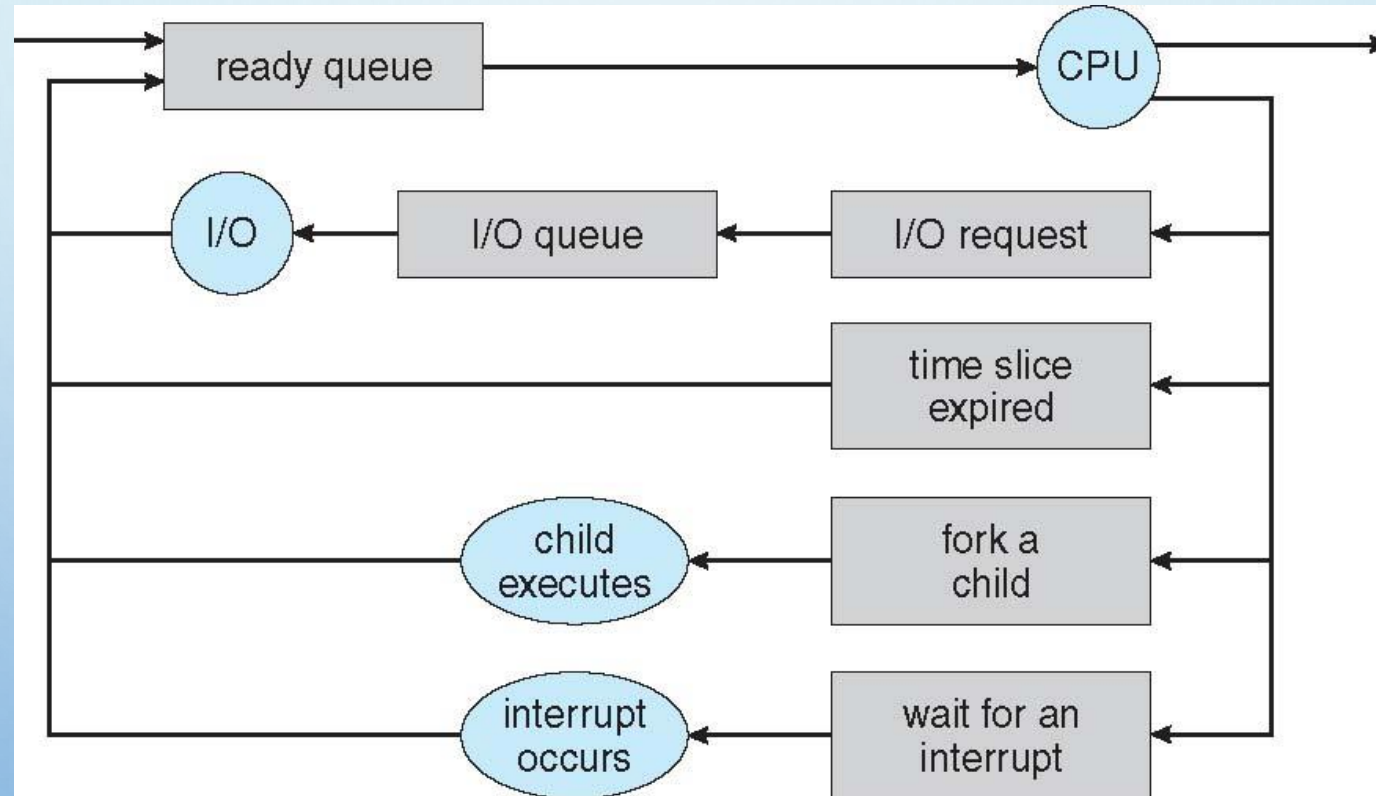
- SO FAR, PROCESS HAS A SINGLE THREAD OF EXECUTION
- CONSIDER HAVING MULTIPLE PROGRAM COUNTERS PER PROCESS
 - MULTIPLE LOCATIONS CAN EXECUTE AT ONCE
 - MULTIPLE THREADS OF CONTROL -> **THREADS**
- MUST THEN HAVE STORAGE FOR THREAD DETAILS, MULTIPLE PROGRAM COUNTERS IN PCB
- SEE NEXT CHAPTER

PROCESS SCHEDULING

- MAXIMIZE CPU USE, QUICKLY SWITCH PROCESSES ONTO CPU FOR TIME SHARING
- **PROCESS SCHEDULER** SELECTS AMONG AVAILABLE PROCESSES FOR NEXT EXECUTION ON CPU
- MAINTAINS **SCHEDULING QUEUES** OF PROCESSES
 - **JOB QUEUE** – SET OF ALL PROCESSES IN THE SYSTEM
 - **READY QUEUE** – SET OF ALL PROCESSES RESIDING IN MAIN MEMORY, READY AND WAITING TO EXECUTE
 - **DEVICE QUEUES** – SET OF PROCESSES WAITING FOR AN I/O DEVICE
 - PROCESSES MIGRATE AMONG THE VARIOUS QUEUES

REPRESENTATION OF PROCESS SCHEDULING

- **Queueing diagram** represents queues, resources, flows



SCHEDULERS

- **SHORT-TERM SCHEDULER** (OR **CPU SCHEDULER**) – SELECTS WHICH PROCESS SHOULD BE EXECUTED NEXT AND ALLOCATES CPU
 - SOMETIMES THE ONLY SCHEDULER IN A SYSTEM
 - SHORT-TERM SCHEDULER IS INVOKED FREQUENTLY (MILLISECONDS) ⇒ (MUST BE FAST)
- **LONG-TERM SCHEDULER** (OR **JOB SCHEDULER**) – SELECTS WHICH PROCESSES SHOULD BE BROUGHT INTO THE READY QUEUE
 - LONG-TERM SCHEDULER IS INVOKED INFREQUENTLY (SECONDS, MINUTES) ⇒ (MAY BE SLOW)
 - THE LONG-TERM SCHEDULER CONTROLS THE **DEGREE OF MULTIPROGRAMMING**
- PROCESSES CAN BE DESCRIBED AS EITHER:
 - **I/O-BOUND PROCESS** – SPENDS MORE TIME DOING I/O THAN COMPUTATIONS, MANY SHORT CPU BURSTS
 - **CPU-BOUND PROCESS** – SPENDS MORE TIME DOING COMPUTATIONS; FEW VERY LONG CPU BURSTS
- LONG-TERM SCHEDULER STRIVES FOR GOOD *PROCESS MIX*

MULTITASKING IN MOBILE SYSTEMS

- SOME MOBILE SYSTEMS (E.G., EARLY VERSION OF IOS) ALLOW ONLY ONE PROCESS TO RUN, OTHERS SUSPENDED
- DUE TO SCREEN REAL ESTATE, USER INTERFACE LIMITS IOS PROVIDES FOR A
 - SINGLE **FOREGROUND** PROCESS– CONTROLLED VIA USER INTERFACE
 - MULTIPLE **BACKGROUND** PROCESSES– IN MEMORY, RUNNING, BUT NOT ON THE DISPLAY, AND WITH LIMITS
 - LIMITS INCLUDE SINGLE, SHORT TASK, RECEIVING NOTIFICATION OF EVENTS, SPECIFIC LONG-RUNNING TASKS LIKE AUDIO PLAYBACK
- ANDROID RUNS FOREGROUND AND BACKGROUND, WITH FEWER LIMITS
 - BACKGROUND PROCESS USES A **SERVICE** TO PERFORM TASKS
 - SERVICE CAN KEEP RUNNING EVEN IF BACKGROUND PROCESS IS SUSPENDED
 - SERVICE HAS NO USER INTERFACE, SMALL MEMORY USE

CONTEXT SWITCH

- WHEN CPU SWITCHES TO ANOTHER PROCESS, THE SYSTEM MUST **SAVE THE STATE** OF THE OLD PROCESS AND LOAD THE **SAVED STATE** FOR THE NEW PROCESS VIA A **CONTEXT SWITCH**
- **CONTEXT** OF A PROCESS REPRESENTED IN THE PCB
- CONTEXT-SWITCH TIME IS OVERHEAD; THE SYSTEM DOES NO USEFUL WORK WHILE SWITCHING
 - THE MORE COMPLEX THE OS AND THE PCB → THE LONGER THE CONTEXT SWITCH
- TIME DEPENDENT ON HARDWARE SUPPORT
 - SOME HARDWARE PROVIDES MULTIPLE SETS OF REGISTERS PER CPU → MULTIPLE CONTEXTS LOADED AT ONCE

OPERATIONS ON PROCESSES

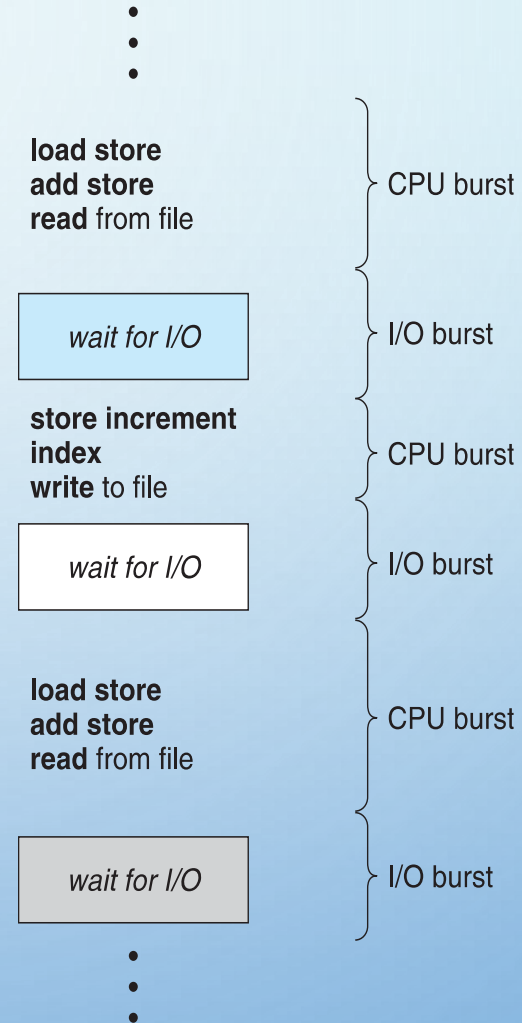
- SYSTEM MUST PROVIDE MECHANISMS FOR:
 - PROCESS CREATION,
 - PROCESS TERMINATION,
 - AND SO ON AS DETAILED NEXT

OBJECTIVES


- TO INTRODUCE CPU SCHEDULING, WHICH IS THE BASIS FOR MULTIPROGRAMMED OPERATING SYSTEMS
- TO DESCRIBE VARIOUS CPU-SCHEDULING ALGORITHMS
- TO DISCUSS EVALUATION CRITERIA FOR SELECTING A CPU-SCHEDULING ALGORITHM FOR A PARTICULAR SYSTEM
- TO EXAMINE THE SCHEDULING ALGORITHMS OF SEVERAL OPERATING SYSTEMS

BASIC CONCEPTS

- MAXIMUM CPU UTILIZATION OBTAINED WITH MULTIPROGRAMMING
- CPU-I/O BURST CYCLE – PROCESS EXECUTION CONSISTS OF A **CYCLE** OF CPU EXECUTION AND I/O WAIT
- **CPU BURST** FOLLOWED BY **I/O BURST**
- CPU BURST DISTRIBUTION IS OF MAIN CONCERN



HISTOGRAM OF CPU-BURST TIMES



CPU SCHEDULER

- **SHORT-TERM SCHEDULER** SELECTS FROM AMONG THE PROCESSES IN READY QUEUE, AND ALLOCATES THE CPU TO ONE OF THEM
 - QUEUE MAY BE ORDERED IN VARIOUS WAYS
- CPU SCHEDULING DECISIONS MAY TAKE PLACE WHEN A PROCESS:
 1. SWITCHES FROM RUNNING TO WAITING STATE
 2. SWITCHES FROM RUNNING TO READY STATE
 3. SWITCHES FROM WAITING TO READY
 4. TERMINATES
- SCHEDULING UNDER 1 AND 4 IS **NONPREEMPTIVE**
- ALL OTHER SCHEDULING IS **PREEMPTIVE**
 - CONSIDER ACCESS TO SHARED DATA
 - CONSIDER PREEMPTION WHILE IN KERNEL MODE
 - CONSIDER INTERRUPTS OCCURRING DURING CRUCIAL OS ACTIVITIES

SCHEDULING CRITERIA

- **CPU UTILIZATION** – KEEP THE CPU AS BUSY AS POSSIBLE
- **THROUGHPUT** – # OF PROCESSES THAT COMPLETE THEIR EXECUTION PER TIME UNIT
- **TURNAROUND TIME** – AMOUNT OF TIME TO EXECUTE A PARTICULAR PROCESS
- **WAITING TIME** – AMOUNT OF TIME A PROCESS HAS BEEN WAITING IN THE READY QUEUE
- **RESPONSE TIME** – AMOUNT OF TIME IT TAKES FROM WHEN A REQUEST WAS SUBMITTED UNTIL THE FIRST RESPONSE IS PRODUCED, NOT OUTPUT (FOR TIME-SHARING ENVIRONMENT)


SCHEDULING ALGORITHM OPTIMIZATION CRITERIA

- MAX CPU UTILIZATION
- MAX THROUGHPUT
- MIN TURNAROUND TIME
- MIN WAITING TIME
- MIN RESPONSE TIME

FIRST- COME, FIRST-SERVED (FCFS) SCHEDULING

<u>PROCESS</u>	<u>BURST TIME</u>
P_1	24
P_2	3
P_3	3

- SUPPOSE THAT THE PROCESSES ARRIVE IN THE ORDER: P_1 , P_2 , P_3
THE GANTT CHART FOR THE SCHEDULE IS:




- WAITING TIME FOR $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- AVERAGE WAITING TIME: $(0 + 24 + 27)/3 = 17$

FCFS SCHEDULING (CONT.)

SUPPOSE THAT THE PROCESSES ARRIVE IN THE ORDER:

$$P_2, P_3, P_1$$

- THE GANTT CHART FOR THE SCHEDULE IS:



- WAITING TIME FOR $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- AVERAGE WAITING TIME: $(6 + 0 + 3)/3 = 3$
- MUCH BETTER THAN PREVIOUS CASE
- **CONVOY EFFECT** – SHORT PROCESS BEHIND LONG PROCESS
 - CONSIDER ONE CPU-BOUND AND MANY I/O-BOUND PROCESSES


SHORTEST-JOB-FIRST (SJF) SCHEDULING

- ASSOCIATE WITH EACH PROCESS THE LENGTH OF ITS NEXT CPU BURST
 - USE THESE LENGTHS TO SCHEDULE THE PROCESS WITH THE SHORTEST TIME
- SJF IS OPTIMAL – GIVES MINIMUM AVERAGE WAITING TIME FOR A GIVEN SET OF PROCESSES
 - THE DIFFICULTY IS KNOWING THE LENGTH OF THE NEXT CPU REQUEST
 - COULD ASK THE USER

EXAMPLE OF SJF

<u>PROCESS</u>	<u>ARRIVA</u>	<u>L TIME</u>	<u>BURST TIME</u>
P_1		0.0	6
P_2		2.0	8
P_3		4.0	7
P_4		5.0	3

- SJF SCHEDULING CHART




- AVERAGE WAITING TIME = $(3 + 16 + 9 + 0) / 4 = 7$

DETERMINING LENGTH OF NEXT CPU BURST

- CAN ONLY ESTIMATE THE LENGTH – SHOULD BE SIMILAR TO THE PREVIOUS ONE
 - THEN PICK PROCESS WITH SHORTEST PREDICTED NEXT CPU BURST
- CAN BE DONE BY USING THE LENGTH OF PREVIOUS CPU BURSTS, USING EXPONENTIAL AVERAGING
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1$
 4. Define : $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.
- COMMONLY, A SET TO $\frac{1}{2}$
- PREEMPTIVE VERSION CALLED **SHORTEST-REMAINING-TIME-FIRST**

PREDICTION OF THE LENGTH OF THE NEXT CPU BURST



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

EXAMPLES OF EXPONENTIAL AVERAGING

- $\alpha = 0$
 - $\tau_{N+1} = \tau_N$
 - RECENT HISTORY DOES NOT COUNT
- $\alpha = 1$
 - $\tau_{N+1} = \alpha T_N$
 - ONLY THE ACTUAL LAST CPU BURST COUNTS

- IF WE EXPAND THE FORMULA, WE GET:

$$\begin{aligned}\tau_{N+1} = & \alpha T_N + (1 - \alpha)\alpha T_{N-1} + \dots \\ & + (1 - \alpha)^j \alpha T_{N-j} + \dots \\ & + (1 - \alpha)^{N+1} \tau_0\end{aligned}$$


- SINCE BOTH α AND $(1 - \alpha)$ ARE LESS THAN OR EQUAL TO 1, EACH SUCCESSIVE TERM HAS LESS WEIGHT THAN ITS PREDECESSOR

EXAMPLE OF SHORTEST-REMAINING-TIME-FIRST

- NOW WE ADD THE CONCEPTS OF VARYING ARRIVAL TIMES AND PREEMPTION TO THE ANALYSIS

PROCESS	ARRIVAL TIME	BURST TIME
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- PREEMPTIVE SJF GANTT CHART



- AVERAGE WAITING TIME = $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5$ MSEC


PRIORITY SCHEDULING

- A PRIORITY NUMBER (INTEGER) IS ASSOCIATED WITH EACH PROCESS
- THE CPU IS ALLOCATED TO THE PROCESS WITH THE HIGHEST PRIORITY (SMALLEST INTEGER \equiv HIGHEST PRIORITY)
 - PREEMPTIVE
 - NONPREEMPTIVE
- SJF IS PRIORITY SCHEDULING WHERE PRIORITY IS THE INVERSE OF PREDICTED NEXT CPU BURST TIME
- PROBLEM \equiv **STARVATION** – LOW PRIORITY PROCESSES MAY NEVER EXECUTE
- SOLUTION \equiv **AGING** – AS TIME PROGRESSES INCREASE THE PRIORITY OF THE PROCESS

EXAMPLE OF PRIORITY SCHEDULING

<u>PROCESS</u>	<u>ARRIVAL TIME</u>	<u>BURST TIME</u>	<u>PRIORITY</u>
P_1		10	3
P_2		1	1
P_3		2	4
P_4		1	5
P_5		5	2

- PRIORITY SCHEDULING GANTT CHART



- AVERAGE WAITING TIME = 8.2 MSEC


ROUND ROBIN (RR)

- EACH PROCESS GETS A SMALL UNIT OF CPU TIME (**TIME QUANTUM** Q), USUALLY 10–100 MILLISECONDS. AFTER THIS TIME HAS ELAPSED, THE PROCESS IS PREEMPTED AND ADDED TO THE END OF THE READY QUEUE.
- IF THERE ARE N PROCESSES IN THE READY QUEUE AND THE TIME QUANTUM IS Q , THEN EACH PROCESS GETS $1/N$ OF THE CPU TIME IN CHUNKS OF AT MOST Q TIME UNITS AT ONCE. NO PROCESS WAITS MORE THAN $(N-1)Q$ TIME UNITS.
- TIMER INTERRUPTS EVERY QUANTUM TO SCHEDULE NEXT PROCESS
- PERFORMANCE
 - Q LARGE \Rightarrow FIFO
 - Q SMALL $\Rightarrow Q$ MUST BE LARGE WITH RESPECT TO CONTEXT SWITCH, OTHERWISE OVERHEAD IS TOO HIGH

EXAMPLE OF RR WITH TIME QUANTUM = 4

<u>PROCESS</u>	<u>BURST TIME</u>
P_1	24
P_2	3
P_3	3

- THE GANTT CHART IS:

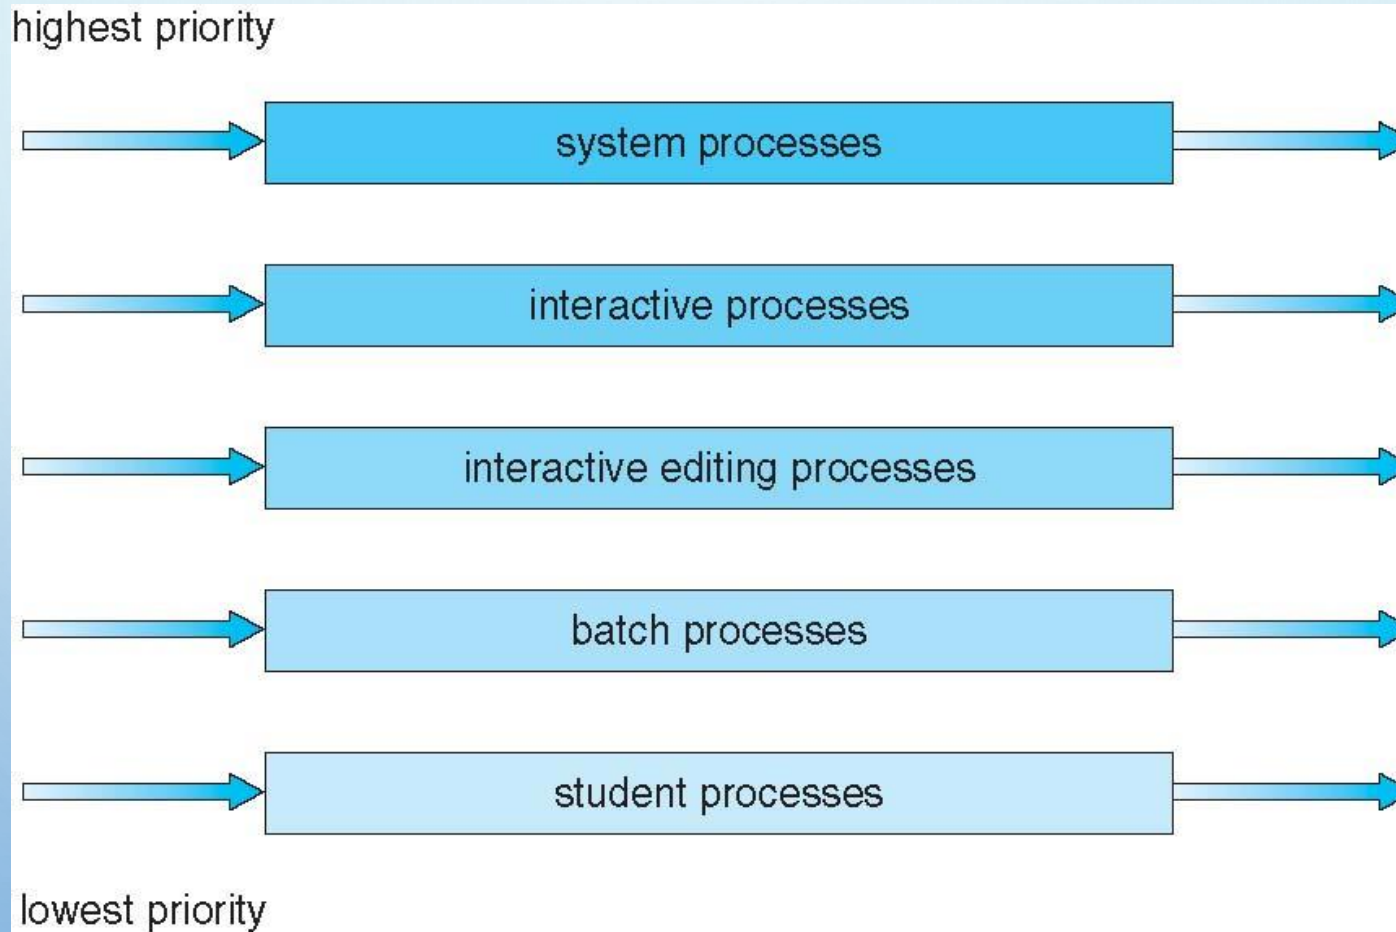


- TYPICALLY, HIGHER AVERAGE TURNAROUND THAN SJF, BUT BETTER *RESPONSE*
- Q SHOULD BE LARGE COMPARED TO CONTEXT SWITCH TIME
- Q USUALLY 10MS TO 100MS, CONTEXT SWITCH < 10 USEC

MULTILEVEL QUEUE

- READY QUEUE IS PARTITIONED INTO SEPARATE QUEUES, EG:
 - **FOREGROUND** (INTERACTIVE)
 - **BACKGROUND** (BATCH)
- PROCESS PERMANENTLY IN A GIVEN QUEUE
- EACH QUEUE HAS ITS OWN SCHEDULING ALGORITHM:
 - FOREGROUND – RR
 - BACKGROUND – FCFS
- SCHEDULING MUST BE DONE BETWEEN THE QUEUES:
 - FIXED PRIORITY SCHEDULING; (I.E., SERVE ALL FROM FOREGROUND THEN FROM BACKGROUND). POSSIBILITY OF STARVATION.
 - TIME SLICE – EACH QUEUE GETS A CERTAIN AMOUNT OF CPU TIME WHICH IT CAN SCHEDULE AMONGST ITS PROCESSES; I.E., 80% TO FOREGROUND IN RR
 - 20% TO BACKGROUND IN FCFS

MULTILEVEL QUEUE SCHEDULING



MULTILEVEL FEEDBACK QUEUE

- A PROCESS CAN MOVE BETWEEN THE VARIOUS QUEUES; AGING CAN BE IMPLEMENTED THIS WAY
- MULTILEVEL-FEEDBACK-QUEUE SCHEDULER DEFINED BY THE FOLLOWING PARAMETERS:
 - NUMBER OF QUEUES
 - SCHEDULING ALGORITHMS FOR EACH QUEUE
 - METHOD USED TO DETERMINE WHEN TO UPGRADE A PROCESS
 - METHOD USED TO DETERMINE WHEN TO DEMOTE A PROCESS
 - METHOD USED TO DETERMINE WHICH QUEUE A PROCESS WILL ENTER WHEN THAT PROCESS NEEDS SERVICE

THREAD SCHEDULING

- DISTINCTION BETWEEN USER-LEVEL AND KERNEL-LEVEL THREADS
- WHEN THREADS SUPPORTED, THREADS SCHEDULED, NOT PROCESSES
- MANY-TO-ONE AND MANY-TO-MANY MODELS, THREAD LIBRARY SCHEDULES USER-LEVEL THREADS TO RUN ON LWP
 - KNOWN AS **PROCESS-CONTENTION SCOPE (PCS)** SINCE SCHEDULING COMPETITION IS WITHIN THE PROCESS
 - TYPICALLY DONE VIA PRIORITY SET BY PROGRAMMER
- KERNEL THREAD SCHEDULED ONTO AVAILABLE CPU IS **SYSTEM-CONTENTION SCOPE (SCS)** – COMPETITION AMONG ALL THREADS IN SYSTEM

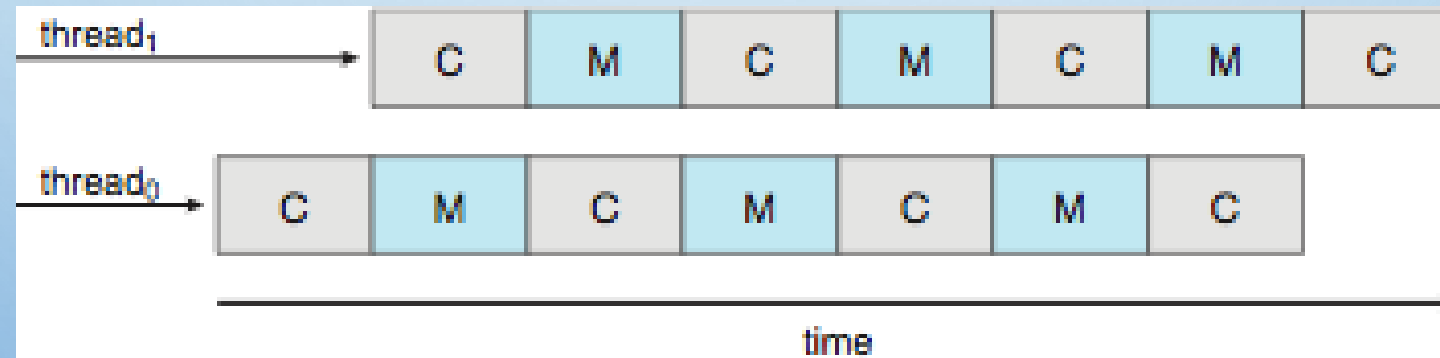
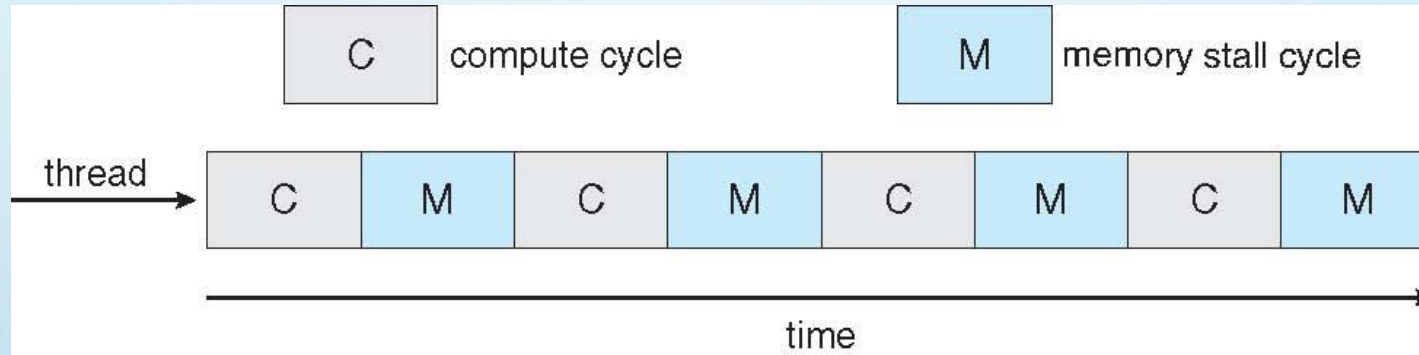
MULTIPLE-PROCESSOR SCHEDULING – LOAD BALANCING

- IF SMP, NEED TO KEEP ALL CPUS LOADED FOR EFFICIENCY
- **LOAD BALANCING** ATTEMPTS TO KEEP WORKLOAD EVENLY DISTRIBUTED
- **PUSH MIGRATION** – PERIODIC TASK CHECKS LOAD ON EACH PROCESSOR, AND IF FOUND PUSHES TASK FROM OVERLOADED CPU TO OTHER CPUS
- **PULL MIGRATION** – IDLE PROCESSORS PULLS WAITING TASK FROM BUSY PROCESSOR

MULTICORE PROCESSORS

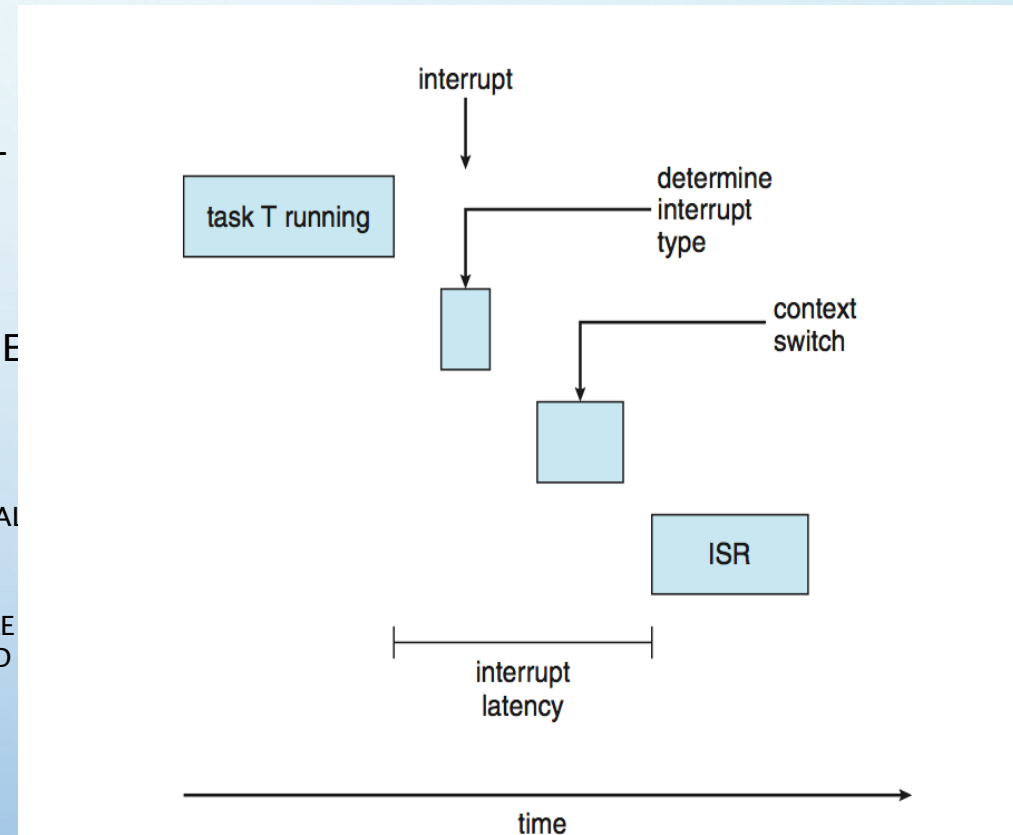
- RECENT TREND TO PLACE MULTIPLE PROCESSOR CORES ON SAME PHYSICAL CHIP
- FASTER AND CONSUMES LESS POWER
- MULTIPLE THREADS PER CORE ALSO GROWING
 - TAKES ADVANTAGE OF MEMORY STALL TO MAKE PROGRESS ON ANOTHER THREAD WHILE MEMORY RETRIEVE HAPPENS

MULTITHREADED MULTICORE SYSTEM



REAL-TIME CPU SCHEDULING

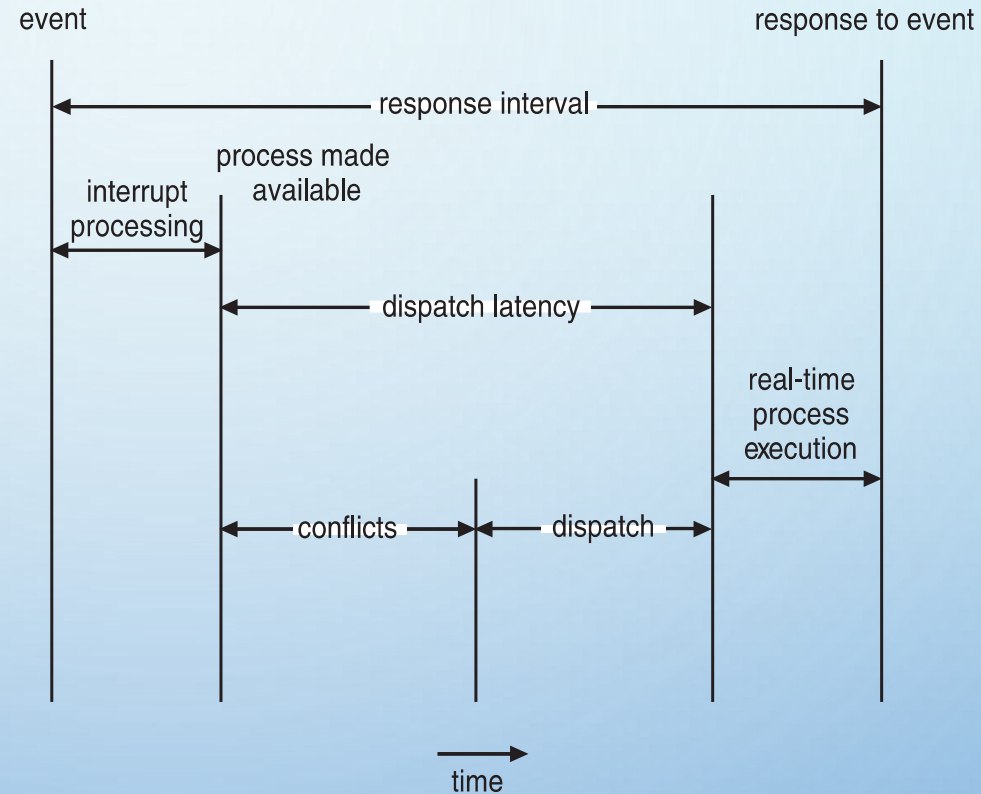
- CAN PRESENT OBVIOUS CHALLENGES
- **SOFT REAL-TIME SYSTEMS** – NO GUARANTEE AS TO WHEN CRITICAL REAL-TIME PROCESS WILL BE SCHEDULED
- **HARD REAL-TIME SYSTEMS** – TASK MUST BE SERVICED BY ITS DEADLINE
- TWO TYPES OF LATENCIES AFFECT PERFORMANCE
 1. INTERRUPT LATENCY – TIME FROM ARRIVAL OF INTERRUPT TO START OF ROUTINE THAT SERVICES INTERRUPT
 2. DISPATCH LATENCY – TIME FOR SCHEDULE TO TAKE CURRENT PROCESS OFF CPU AND SWITCH TO ANOTHER



REAL-TIME CPU SCHEDULING (CONT.)

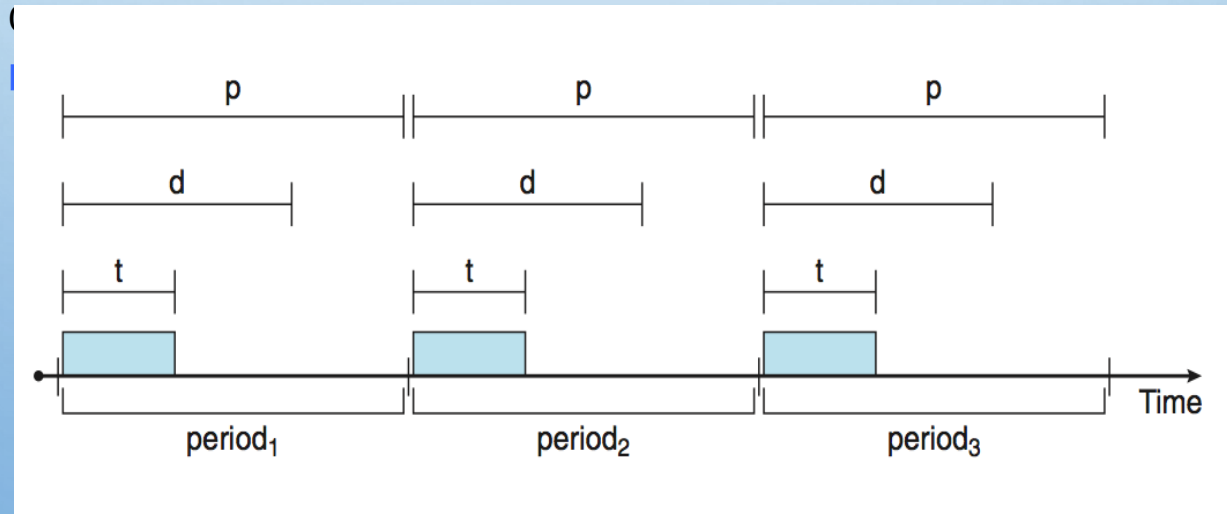
- CONFLICT PHASE OF DISPATCH LATENCY:

1. PREEMPTION OF ANY PROCESS RUNNING IN KERNEL MODE
2. RELEASE BY LOW-PRIORITY PROCESS OF RESOURCES NEEDED BY HIGH-PRIORITY PROCESSES



PRIORITY-BASED SCHEDULING

- FOR REAL-TIME SCHEDULING, SCHEDULER MUST SUPPORT PREEMPTIVE, PRIORITY-BASED SCHEDULING
 - BUT ONLY GUARANTEES SOFT REAL-TIME
- FOR HARD REAL-TIME MUST ALSO PROVIDE ABILITY TO MEET DEADLINES
- PROCESSES HAVE NEW CHARACTERISTICS: **PERIODIC** ONES REQUIRE CPU AT CONSTANT INTERVALS
 - HAS PROCESSING TIME T , DEADLINE D , PERIOD P
 - (
 - I




VIRTUALIZATION AND SCHEDULING


- VIRTUALIZATION SOFTWARE SCHEDULES MULTIPLE GUESTS ONTO CPU(S)
- EACH GUEST DOING ITS OWN SCHEDULING
 - NOT KNOWING IT DOESN'T OWN THE CPUS
 - CAN RESULT IN POOR RESPONSE TIME
 - CAN EFFECT TIME-OF-DAY CLOCKS IN GUESTS
- CAN UNDO GOOD SCHEDULING ALGORITHM EFFORTS OF GUESTS

RATE MONOTONIC SCHEDULING

- A PRIORITY IS ASSIGNED BASED ON THE INVERSE OF ITS PERIOD
- SHORTER PERIODS = HIGHER PRIORITY;
- LONGER PERIODS = LOWER PRIORITY



MISSED DEADLINES WITH RATE MONOTONIC SCHEDULING




EARLIEST DEADLINE FIRST SCHEDULING (EDF)

- PRIORITIES ARE ASSIGNED ACCORDING TO DEADLINES:

THE EARLIER THE DEADLINE, THE HIGHER THE PRIORITY;

THE LATER THE DEADLINE, THE LOWER THE PRIORITY



PROPORTIONAL SHARE SCHEDULING

- T SHARES ARE ALLOCATED AMONG ALL PROCESSES IN THE SYSTEM
- AN APPLICATION RECEIVES N SHARES WHERE $N < T$
- THIS ENSURES EACH APPLICATION WILL RECEIVE N / T OF THE TOTAL PROCESSOR TIME

OPERATING SYSTEM EXAMPLES

- LINUX SCHEDULING
- WINDOWS SCHEDULING
- SOLARIS SCHEDULING